

Copy of CAD_DLProject_Challenge-1_AlexNet_V12.ipynb

File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

[1] from google.colab import drive
drive.mount('/content/gdrive', force_remount=True)

Mounted at /content/gdrive

[x] pip install torch-lr-finder

Collecting torch-lr-finder

 Downloading torch_lr_finder-0.2.1-py3-none-any.whl (11 kB)

Requirement already satisfied: numpy >=1.16.0 in /usr/local/lib/python3.7/dist-packages (from torch-lr-finder) (1.19.5)

Requirement already satisfied: torch<1.0.0,!=1.0.0rc1 in /usr/local/lib/python3.7/dist-packages (from torch-lr-finder) (3.2.2)

Requirement already satisfied: torch<0.4.1 in /usr/local/lib/python3.7/dist-packages (from torch-lr-finder) (1.10.0+cull11)

Requirement already satisfied: packaging in /usr/local/lib/python2.7/dist-packages (from torch-lr-finder) (21.3)

Requirement already satisfied: tqdm in /usr/local/lib/python3.7/dist-packages (from torch-lr-finder) (4.62.3)

Requirement already satisfied: typing-extensions in /usr/local/lib/python3.7/dist-packages (from torch<0.4.1>torch-lr-finder) (3.10.0.2)

Requirement already satisfied: pyarsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in /usr/local/lib/python3.7/dist-packages (from matplotlib>torch-lr-finder) (3.0.6)

Requirement already satisfied: pyparsing<2.1.2,!=2.1.6,>=2.0.1 in /usr/local/lib/python3.7/dist-packages (from torch-lr-finder) (0.11.0)

Requirement already satisfied: python-dateutil<2.1 in /usr/local/lib/python3.7/dist-packages (from matplotlib>torch-lr-finder) (2.8.2)

Requirement already satisfied: six=>1.5 in /usr/local/lib/python3.7/dist-packages (from python-dateutil=>2.1>matplotlib>torch-lr-finder) (1.15.0)

Installing collected packages: torch-lr-finder

Successfully installed torch-lr-finder-0.2.1

[2] pip install torch-lr-finder -v --global-option="apex"

Value for scheme.purelib does not match. Please report this to <<https://github.com/pypa/pip/issues/9617>>

distutils: /usr/local/lib/python3.7/dist-packages

sysconfig: /usr/lib/python3.7/_sitecustomize.py

Value for scheme.headers does not match. Please report this to <<https://github.com/pypa/pip/issues/9617>>

distutils: /usr/local/include/python3.7/UNKNOWN

sysconfig: /usr/include/python3.7m/UNKNOWN

Value for scheme.scripts does not match. Please report this to <<https://github.com/pypa/pip/issues/9617>>

distutils: /usr/local/bin

sysconfig: /usr/bin

Value for scheme.data does not match. Please report this to <<https://github.com/pypa/pip/issues/9617>>

distutils: /usr/local

sysconfig: /usr

Additional context:

user = False

home = None

root = None

prefix = None

Non-user install because site-packages writeable

Created temporary directory: /tmp/pip-ephem-wheel-cache-yakor88j

Created temporary directory: /tmp/pip-req-tracker-fbftwcis

Initialized build tracking at /tmp/pip-req-tracker-fbftwcis

Created temporary directory: /tmp/pip-req-tracker-fbftwcis

Entered build tracker: /tmp/pip-req-tracker-fbftwcis

Created temporary directory: /tmp/pip-install-x90j94q

Requirement already satisfied: torch-lr-finder in /usr/local/lib/python3.7/dist-packages (0.2.1)

Requirement already satisfied: numpy in /usr/local/lib/python3.7/dist-packages (from torch-lr-finder) (1.19.5)

Requirement already satisfied: tqdm in /usr/local/lib/python3.7/dist-packages (from torch-lr-finder) (4.62.3)

Requirement already satisfied: packaging in /usr/local/lib/python3.7/dist-packages (from torch-lr-finder) (1.10.0+cull11)

Requirement already satisfied: torch<0.4.1 in /usr/local/lib/python3.7/dist-packages (from torch-lr-finder) (1.10.0+cull11)

Requirement already satisfied: pyParsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in /usr/local/lib/python3.7/dist-packages (from torch-lr-finder) (3.10.0.2)

Requirement already satisfied: pyparsing<2.1.2,!=2.1.6,>=2.0.1 in /usr/local/lib/python3.7/dist-packages (from torch-lr-finder) (0.11.0)

Requirement already satisfied: python-dateutil<2.1 in /usr/local/lib/python3.7/dist-packages (from matplotlib>torch-lr-finder) (2.8.2)

Requirement already satisfied: cyclene<0.1.0,!=0.1.0rc1 in /usr/local/lib/python3.7/dist-packages (from matplotlib>torch-lr-finder) (0.1.0)

Requirement already satisfied: six=>1.5 in /usr/local/lib/python3.7/dist-packages (from python-dateutil=>2.1>matplotlib>torch-lr-finder) (1.15.0)

Created temporary directory: /tmp/pip-unpack-v8spfdl

Value for scheme.distlib does not match. Please report this to <<https://github.com/pypa/pip/issues/9617>>

distutils: /usr/local/lib/python3.7/dist-packages

sysconfig: /usr/lib/python3.7/_sitecustomize.py

Value for scheme.headers does not match. Please report this to <<https://github.com/pypa/pip/issues/9617>>

distutils: /usr/local/lib/python3.7/dist-packages

sysconfig: /usr/lib/python3.7/_sitecustomize.py

Value for scheme.scripts does not match. Please report this to <<https://github.com/pypa/pip/issues/9617>>

distutils: /usr/local/include/python3.7/UNKNOWN

sysconfig: /usr/include/python3.7m/UNKNOWN

Value for scheme.scripts does not match. Please report this to <<https://github.com/pypa/pip/issues/9617>>

distutils: /usr/local/bin

sysconfig: /usr/bin

Value for scheme.data does not match. Please report this to <<https://github.com/pypa/pip/issues/9617>>

distutils: /usr/local

sysconfig: /usr

Additional context:

user = False

home = None

root = None

prefix = None

Removed build tracker: '/tmp/pip-req-tracker-fbftwcis'

[3]

[4] import torch, torchvision
from torchvision import datasets, models, transforms
import torch.nn as nn
import torch.optim as optim
from torch.optim import lr_scheduler
from torch.utils.data import DataLoader
import time
from torchsummary import summary

import numpy as np
import matplotlib.pyplot as plt
import os

from PIL import Image
from torch_lr_finder import LRFinder

[5] if torch.cuda.is_available():
 device = torch.device("cuda:0")
 print("GPU")
 torch.cuda.device_count()
 torch.cuda.get_device_name(0)
 !nvidia-smi
else:
 device = torch.device("cpu")

GPU

Mon Jan 10 13:05:41 2022

NVIDIA-SMI 495.44	Driver Version: 460.32.03	CUDA Version: 11.2		
+-----+-----+-----+				
GPU Name	Persistence-M	Bus-Id	Disp.A	Volatile Uncorr. ECC
Fan Temp Perf Pwr/Usage/Cap			Memory-Usage	GPU-Util Compute M.
				MIG M.
+-----+-----+-----+				
0 Tesla P100-PCIE.. Off 00000000:00:04.0 Off	0			
N/A 34C P0 26W / 250W 2MLB / 16280M1B 0% Default				
+-----+-----+-----+				
Processes:				
GPU GI CI PID Type Process name			GPU Memory	
ID ID ID			Usage	
+-----+-----+-----+				
No running processes found				

[6] # parameters
Batch_Size = 8
Number_of_Epoch = 30
Init_LR_for_search = 1e-7

[7] # Applying Transforms to the Data

```
image_transforms = [  
    'train': transforms.Compose([  
        transforms.RandomResizedCrop(size = 400, scale=(0.8, 1.0)), #size=256  
        transforms.RandomRotation(degrees=15),  
        transforms.RandomHorizontalFlip(),  
        transforms.CenterCrop(size = 300), #size=224  
        transforms.ToTensor(),  
        transforms.Normalize([0.485, 0.456, 0.406],  
                         [0.229, 0.224, 0.225])
```

```

    ''',
    'valid': transforms.Compose([
        transforms.Resize(size = 400), #size=256
        transforms.CenterCrop(size = 300), #size=224
        transforms.ToTensor(),
        transforms.Normalize([0.485, 0.456, 0.406],
                           [0.229, 0.224, 0.225])
    ]),
    'test': transforms.Compose([
        transforms.Resize(size = 400), #size=256
        transforms.CenterCrop(size = 300), #size=224
        transforms.ToTensor(),
        transforms.Normalize([0.485, 0.456, 0.406],
                           [0.229, 0.224, 0.225])
    ])
}

[8] # Load the Data
# Set train and valid directory paths
dataset = '/content/gdrive/MyDrive/cl'
#train_path = ("/content/gdrive/MyDrive/CAD Project Data/challenge1/train")
#val_path = ("/content/gdrive/MyDrive/CAD Project Data/challenge1/val")

train_directory = os.path.join(dataset, 'train')
valid_directory = os.path.join(dataset, 'val')

# Batch size
bs = Batch Size

# Number of classes
num_classes = len(os.listdir(valid_directory)) #10#2#257
print(num_classes)

# Load Data from folders
data = {
    'train': datasets.ImageFolder(root=train_directory, transform=image_transforms['train']),
    'val': datasets.ImageFolder(root=valid_directory, transform=image_transforms['valid'])
}

# Get a mapping of the indices to the class names, in order to see the output classes of the test images.
idx_to_class = {v: k for k, v in data['train'].class_to_idx.items()}
print(idx_to_class)

# Size of Data, to be used for calculating Average Loss and Accuracy
train_data_size = len(data['train'])
valid_data_size = len(data['val'])

# Create iterators for the Data loaded using DataLoader module
train_data_loader = DataLoader(data['train'], batch_size=bs, shuffle=True)
valid_data_loader = DataLoader(data['val'], batch_size=bs, shuffle=True)

2
{0: 'les', 1: 'nv'}
```

[9] train_data_size, valid_data_size

```
(4800, 1200)
```

[10] model = models.alexnet(pretrained=True)
model.to(device)
model

Downloading: "https://download.pytorch.org/models/alexnet-owt-7be5be79.pth" to /root/.cache/torch/hub/checkpoints/alexnet-owt-7be5be79.pth

100% [██████████] 233M/233M [00:01<00:00, 233MB/s]

AlexNet(
 (features): Sequential(
 (0): Conv2d(3, 64, kernel_size=(11, 11), stride=(4, 4), padding=(2, 2))
 (1): ReLU(inplace=True)
 (2): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1, ceil_mode=False)
 (3): Conv2d(64, 192, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
 (4): ReLU(inplace=True)
 (5): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1, ceil_mode=False)
 (6): Conv2d(192, 384, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
 (7): ReLU(inplace=True)
 (8): Conv2d(384, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
 (9): ReLU(inplace=True)
 (10): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
 (11): ReLU(inplace=True)
 (12): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1, ceil_mode=False)
) (avgpool): AdaptiveAvgPool2d(output_size=(6, 6))
 (classifier): Sequential(
 (0): Dropout(p=0.5, inplace=False)
 (1): Linear(in_features=9216, out_features=4096, bias=True)
 (2): ReLU(inplace=True)
 (3): Dropout(p=0.5, inplace=False)
 (4): Linear(in_features=4096, out_features=4096, bias=True)
 (5): ReLU(inplace=True)
 (6): Linear(in_features=4096, out_features=1000, bias=True)
)
)

[14] # Freeze model parameters
for param in model.parameters():
 param.requires_grad = False

[15] # Change the final layer of AlexNet Model for Transfer Learning
model.classifier[6] = nn.Linear(4096, num_classes)
model.classifier.add_module("7", nn.LogSoftmax(dim = 1))
model.to(device)
model

AlexNet(
 (features): Sequential(
 (0): Conv2d(3, 64, kernel_size=(11, 11), stride=(4, 4), padding=(2, 2))
 (1): ReLU(inplace=True)
 (2): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1, ceil_mode=False)
 (3): Conv2d(64, 192, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
 (4): ReLU(inplace=True)
 (5): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1, ceil_mode=False)
 (6): Conv2d(192, 384, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
 (7): ReLU(inplace=True)
 (8): Conv2d(384, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
 (9): ReLU(inplace=True)
 (10): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
 (11): ReLU(inplace=True)
 (12): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1, ceil_mode=False)
) (avgpool): AdaptiveAvgPool2d(output_size=(6, 6))
 (classifier): Sequential(
 (0): Dropout(p=0.5, inplace=False)
 (1): Linear(in_features=9216, out_features=4096, bias=True)
 (2): ReLU(inplace=True)
 (3): Dropout(p=0.5, inplace=False)
 (4): Linear(in_features=4096, out_features=4096, bias=True)
 (5): ReLU(inplace=True)
 (6): Linear(in_features=4096, out_features=2, bias=True)
 (7): LogSoftmax(dim=1)
)
)

[16] summary(model, (3, 224, 224))

Layer (type)	Output Shape	Param #
Conv2d-1	[-, 64, 55, 55]	23,296
ReLU-2	[-, 64, 55, 55]	0
MaxPool2d-3	[-, 64, 27, 27]	0
Conv2d-4	[-, 192, 27, 27]	307,392
ReLU-5	[-, 192, 27, 27]	0
MaxPool2d-6	[-, 192, 13, 13]	0
Conv2d-7	[-, 384, 13, 13]	663,936
ReLU-8	[-, 384, 13, 13]	0
Conv2d-9	[-, 256, 13, 13]	884,992
ReLU-10	[-, 256, 13, 13]	0
Conv2d-11	[-, 256, 13, 13]	590,080
ReLU-12	[-, 256, 13, 13]	0
MaxPool2d-13	[-, 256, 6, 6]	0
AdaptiveAvgPool2d-14	[-, 256, 6, 6]	0
Dropout-15	[-, 9216]	0
Linear-16	[-, 4096]	37,752,832
ReLU-17	[-, 4096]	0
Drotnout-18	[-, 4096]	0

```

--> Linear-19      [-1, 4096]    16,781,312
  ReLU-20        [-1, 4096]      0
  Linear-21      [-1, 2]        8,194
  LogSoftmax-22   [-1, -1]      0
=====
Total params: 57,012,034
Trainable params: 8,194
Non-trainable params: 57,003,840
-----
Input size (MB): 0.57
Forward/backward pass size (MB): 8.37
Params size (MB): 217.48
Estimated Total Size (MB): 226.43
-----
```

```

[17] # Define Optimizer and Loss Function
loss_func = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr = Init_LR_for_search, weight_decay = 1e-2)
lr_finder = LRFinder(model, optimizer, loss_func, device = "cuda:0")
lr_finder.range_test(train_data_loader, end_lr = 100, num_iter = 100)
lr_finder.plot()
#optimizer
```

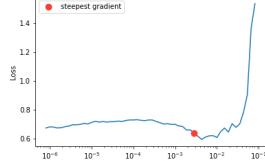
70% 70/100 [08:29<03:16, 6.56s/it]

Stopping early, the loss has diverged

Learning rate search finished. See the graph with {finder_name}.plot()

LR suggestion: steepest gradient

Suggested LR: 2.85E-03



(<matplotlib.axes._subplots.AxesSubplot at 0x7fa72a77c090>, 0.002848035868435803)

```
[18] optimizer = optim.Adam(model.parameters(), lr = 2.85E-03)
```

```
[19] decaying_lr = lr_scheduler.StepLR(optimizer, step_size = 4, gamma = 0.1)
```

```

def train_and_validate(model, loss_criterion, optimizer, scheduler, epochs):
    ...
    Function to train and validate
    Parameters
    :param model: Model to train and validate
    :param loss_criterion: Loss Criterion to minimize
    :param optimizer: Optimizer for computing gradients
    :param epochs: Number of epochs (default=25)

    Returns
    :model: Trained Model with best validation accuracy
    :history: (dict object): Having training loss, accuracy and validation loss, accuracy
    ...

    start = time.time()
    history = []
    best_acc = 0.0

    for epoch in range(epochs):
        epoch_start = time.time()
        print("Epoch: {}/{}".format(epoch+1, epochs))

        # Set to training mode
        model.train()

        # Loss and Accuracy within the epoch
        train_loss = 0.0
        train_acc = 0.0

        valid_loss = 0.0
        valid_acc = 0.0

        for i, (inputs, labels) in enumerate(train_data_loader):
            inputs = inputs.to(device)
            labels = labels.to(device)

            # Clean existing gradients
            optimizer.zero_grad()

            # Forward pass - compute outputs on input data using the model
            outputs = model(inputs).to(device)

            # Compute loss
            loss = loss_criterion(outputs, labels)

            # Backpropagate the gradients
            loss.backward()

            # Update the parameters
            optimizer.step()

            # Compute the total loss for the batch and add it to train_loss
            train_loss += loss.item() * inputs.size(0)

            # Compute the accuracy
            ret, predictions = torch.max(outputs.data, 1)
            correct_counts = predictions.eq(labels.data.view_as(predictions))

            # Convert correct_counts to float and then compute the mean
            acc = torch.mean(correct_counts.type(torch.FloatTensor))

            # Compute total accuracy in the whole batch and add to train_acc
            train_acc += acc.item() * inputs.size(0)

            #print("Batch number: {}3d, Training: Loss: {:.4f}, Accuracy: {:.4f}".format(i, loss.item(), acc.item()))
        scheduler.step()

        # Validation - No gradient tracking needed
        with torch.no_grad():

            # Set to evaluation mode
            model.eval()

            # Validation loop
            for j, (inputs, labels) in enumerate(valid_data_loader):
                inputs = inputs.to(device)
                labels = labels.to(device)

                # Forward pass - compute outputs on input data using the model
                outputs = model(inputs).to(device)

                # Compute loss
                loss = loss_criterion(outputs, labels).to(device)

                # Compute the total loss for the batch and add it to valid_loss
                valid_loss += loss.item() * inputs.size(0)

                # Calculate validation accuracy
                ret, predictions = torch.max(outputs.data, 1)
                correct_counts = predictions.eq(labels.data.view_as(predictions))

                # Convert correct_counts to float and then compute the mean
                acc = torch.mean(correct_counts.type(torch.FloatTensor))

                # Compute total accuracy in the whole batch and add to valid_acc
                valid_acc += acc.item() * inputs.size(0)

                #print("Validation Batch number: {}3d, Validation: Loss: {:.4f}, Accuracy: {:.4f}".format(j, loss.item(), acc.item()))

            # Find average training loss and training accuracy
            avg_train_loss = train_loss/train_data_size
```

```

avg_train_acc = train_acc/train_data_size

# Find average training loss and training accuracy
avg_valid_loss = valid_loss/valid_data_size
avg_valid_acc = valid_acc/valid_data_size

history.append([avg_train_loss, avg_valid_loss, avg_train_acc, avg_valid_acc])

epoch_end = time.time()

print("Epoch : {:03d}, Training: Loss: {:.4f}, Accuracy: {:.4f}%, Validation : Loss : {:.4f}, Accuracy: {:.4f}%, Time: {:.4f}s".format(epoch+1, avg_train_loss, avg_train_acc*100, avg_valid_loss, avg_valid_acc, epoch_end))

# Save if the model has best accuracy till now
torch.save(model, dataset+'model'+str(epoch)+'.pt')

return model, history

```

```

num_epochs = Number_of_Epoch
trained_model, history = train_and_validate(model, loss_func, optimizer, decaying_lr, num_epochs)

torch.save(history, dataset+'history.pt')

```

```

history = np.array(history)
plt.plot(history[:,0:2])
plt.legend(['Tr Loss', 'Val Loss'])
plt.xlabel('Epoch Number')
plt.ylabel('Loss')
plt.ylim(0,1)
plt.savefig(dataset+'_loss_curve.png')
plt.show()

```

```

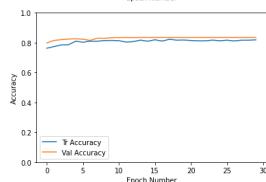
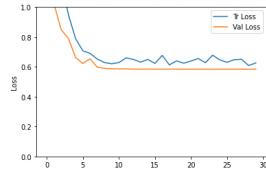
plt.plot(history[:,2:4])
plt.legend(['Tr Accuracy', 'Val Accuracy'])
plt.xlabel('Epoch Number')
plt.ylabel('Accuracy')
plt.ylim(0,1)
plt.savefig(dataset+'_accuracy_curve.png')
plt.show()

```

```

Epoch : 022, Training: Loss: 0.6556, Accuracy: 81.2292%, Validation : Loss : 0.5846, Accuracy: 83.4167%, Time: 92.5055s
Epoch: 23/30
Epoch : 023, Training: Loss: 0.6277, Accuracy: 81.2292%, Validation : Loss : 0.5846, Accuracy: 83.4167%, Time: 92.2805s
Epoch: 24/30
Epoch : 024, Training: Loss: 0.6781, Accuracy: 81.7083%, Validation : Loss : 0.5846, Accuracy: 83.4167%, Time: 92.9448s
Epoch: 25/30
Epoch : 025, Training: Loss: 0.6466, Accuracy: 81.2788%, Validation : Loss : 0.5846, Accuracy: 83.4167%, Time: 92.8627s
Epoch: 26/30
Epoch : 026, Training: Loss: 0.6303, Accuracy: 81.6667%, Validation : Loss : 0.5846, Accuracy: 83.4167%, Time: 92.8035s
Epoch: 27/30
Epoch : 027, Training: Loss: 0.6479, Accuracy: 81.1458%, Validation : Loss : 0.5846, Accuracy: 83.4167%, Time: 92.8874s
Epoch: 28/30
Epoch : 028, Training: Loss: 0.6508, Accuracy: 81.6667%, Validation : Loss : 0.5846, Accuracy: 83.4167%, Time: 92.7168s
Epoch: 29/30
Epoch : 029, Training: Loss: 0.6089, Accuracy: 81.6667%, Validation : Loss : 0.5846, Accuracy: 83.4167%, Time: 92.6678s
Epoch: 30/30
Epoch : 030, Training: Loss: 0.6263, Accuracy: 81.8758%, Validation : Loss : 0.5846, Accuracy: 83.4167%, Time: 93.4204s

```



```

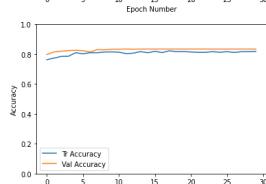
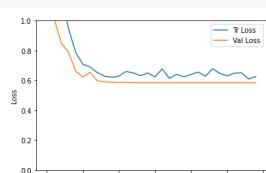
[22] history = np.array(history)
plt.plot(history[:,0:2])
plt.legend(['Tr Loss', 'Val Loss'])
plt.xlabel('Epoch Number')
plt.ylabel('Loss')
plt.ylim(0,1)
plt.savefig(dataset+'_loss_curve.png')
plt.show()

```

```

plt.plot(history[:,2:4])
plt.legend(['Tr Accuracy', 'Val Accuracy'])
plt.xlabel('Epoch Number')
plt.ylabel('Accuracy')
plt.ylim(0,1)
plt.savefig(dataset+'_accuracy_curve.png')
plt.show()

```



```
[ ]
```

```

[ ] def predict(model, test_image_name):
    ...
    Function to predict the class of a single test image
    Parameters
    :param model: Model to test
    :param test_image_name: Test image
    ...

```

```
transform = image_transforms['test']

test_image = Image.open(test_image_name)
plt.imshow(test_image)

test_image_tensor = transform(test_image)

if torch.cuda.is_available():
    test_image_tensor = test_image_tensor.view(1, 3, 224, 224).cuda()
else:
    test_image_tensor = test_image_tensor.view(1, 3, 224, 224)

with torch.no_grad():
    model.eval()
    # Model outputs log probabilities
    out = model(test_image_tensor)
    ps = torch.exp(out)
    topk, topclass = ps.topk(3, dim=1)
    for i in range(3):
        print("Prediction", i+1, ":", idx_to_class[topclass.numpy()[0][i]], ", Score: ", topk.numpy()[0][i])

[ ] # predict(trained_model, 'caltec256subset/test/009_0008.jpg')
# predict(trained_model, 'caltec256subset/test/090_0107.jpg')
# predict(trained_model, 'caltec256subset/test/150_0044.jpg')
```

✓ 0s completed at 4:11 PM

● X