

Copy of CAD\_DLProject\_Challenge-2\_ResNet\_V12.ipynb

File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

[1] from google.colab import drive  
drive.mount('/content/gdrive', force\_remount=True)

Mounted at /content/gdrive

[x] [2] pip install torch-lr-finder

Collecting torch-lr-finder  
 Downloading torch\_lr\_finder-0.2.1-py3-none-any.whl (11 kB)  
Requirement already satisfied: packaging in /usr/local/lib/python3.7/dist-packages (from torch-lr-finder) (0.1.3)  
Requirement already satisfied: requests in /usr/local/lib/python3.7/dist-packages (from torch-lr-finder) (2.22.2)  
Requirement already satisfied: numpy in /usr/local/lib/python3.7/dist-packages (from torch-lr-finder) (1.19.5)  
Requirement already satisfied: torch<=0.4.1 in /usr/local/lib/python3.7/dist-packages (from torch-lr-finder) (0.4.2.3)  
Requirement already satisfied: tqdm in /usr/local/lib/python3.7/dist-packages (from torch-lr-finder) (4.62.3)  
Requirement already satisfied: typing-extensions in /usr/local/lib/python3.7/dist-packages (from torch<=0.4.1->torch-lr-finder) (3.10.0.2)  
Requirement already satisfied: pyrsparser!=2.0.4,!>=2.1.6,!>=2.0.1 in /usr/local/lib/python3.7/dist-packages (from matplotlib->torch-lr-finder) (3.0.6)  
Requirement already satisfied: python-dateutil>=2.1 in /usr/local/lib/python3.7/dist-packages (from matplotlib->torch-lr-finder) (2.8.2)  
Requirement already satisfied: six<=1.15.0 in /usr/local/lib/python3.7/dist-packages (from matplotlib->torch-lr-finder) (1.13.2)  
Requirement already satisfied: colorsys>=0.10 in /usr/local/lib/python3.7/dist-packages (from matplotlib->torch-lr-finder) (0.11.0)  
Requirement already satisfied: six=>1.5 in /usr/local/lib/python3.7/dist-packages (from python-dateutil>=2.1->matplotlib->torch-lr-finder) (1.15.0)  
Installing collected packages: torch-lr-finder  
Successfully installed torch-lr-finder-0.2.1

[x] [3] pip install torch-lr-finder -v --global-option='apex'

Value for scheme.purelib does not match. Please report this to <<https://github.com/pypa/pip/issues/9617>>  
distutils: /usr/local/lib/python3.7/dist-packages  
sysconfig: /usr/lib/python3.7/\_sitecustomize.py  
Value for scheme.headers does not match. Please report this to <<https://github.com/pypa/pip/issues/9617>>  
distutils: /usr/local/include/python3.7/UNKNOWN  
sysconfig: /usr/include/python3.7m/UNKNOWN  
Value for scheme.scripts does not match. Please report this to <<https://github.com/pypa/pip/issues/9617>>  
distutils: /usr/local/bin  
sysconfig: /usr/bin  
Value for scheme.data does not match. Please report this to <<https://github.com/pypa/pip/issues/9617>>  
distutils: /usr/local  
sysconfig: /usr  
Additional context:  
user = False  
home = None  
root = None  
prefix = None  
Non-user install because site-packages writeable  
Created temporary directory: /tmp/pip-ephem-wheel-cache-qqbtb3q2  
Created temporary directory: /tmp/pip-req-tracker-3\_a003f6  
Initialized build tracking at /tmp/pip-req-tracker-3\_a003f6  
Created temporary directory: /tmp/pip-req-tracker-3\_a003f6  
Entered build tracker: /tmp/pip-req-tracker-3\_a003f6  
Created temporary directory: /tmp/pip-install-slyvym4s  
Requirement already satisfied: torch-lr-finder in /usr/local/lib/python3.7/dist-packages (0.2.1)  
Requirement already satisfied: packaging in /usr/local/lib/python3.7/dist-packages (from torch-lr-finder) (21.3)  
Requirement already satisfied: tqdm in /usr/local/lib/python3.7/dist-packages (from torch-lr-finder) (4.62.3)  
Requirement already satisfied: torch<=0.4.1 in /usr/local/lib/python3.7/dist-packages (from torch-lr-finder) (0.4.2.3)  
Requirement already satisfied: numpy in /usr/local/lib/python3.7/dist-packages (from torch-lr-finder) (1.19.5)  
Requirement already satisfied: typing-extensions in /usr/local/lib/python3.7/dist-packages (from torch<=0.4.1->torch-lr-finder) (3.10.0.2)  
Requirement already satisfied: pyrsparser>=1.0.1 in /usr/local/lib/python3.7/dist-packages (from matplotlib->torch-lr-finder) (1.13.2)  
Requirement already satisfied: python-dateutil>=2.1 in /usr/local/lib/python3.7/dist-packages (from matplotlib->torch-lr-finder) (2.8.2)  
Requirement already satisfied: colorsys>=0.10 in /usr/local/lib/python3.7/dist-packages (from matplotlib->torch-lr-finder) (0.11.0)  
Requirement already satisfied: six=>1.5 in /usr/local/lib/python3.7/dist-packages (from python-dateutil>=2.1->matplotlib->torch-lr-finder) (1.15.0)  
Created temporary directory: /tmp/pip-unpack-dvv66lm  
Value for scheme.platlib does not match. Please report this to <<https://github.com/pypa/pip/issues/9617>>  
distutils: /usr/local/lib/python3.7/dist-packages  
sysconfig: /usr/lib/python3.7/\_site\_packages  
Value for scheme.headers does not match. Please report this to <<https://github.com/pypa/pip/issues/9617>>  
distutils: /usr/local/lib/python3.7/dist-packages  
sysconfig: /usr/lib/python3.7/\_site\_packages  
Value for scheme.scripts does not match. Please report this to <<https://github.com/pypa/pip/issues/9617>>  
distutils: /usr/local/include/python3.7/UNKNOWN  
sysconfig: /usr/include/python3.7m/UNKNOWN  
Value for scheme.scripts does not match. Please report this to <<https://github.com/pypa/pip/issues/9617>>  
distutils: /usr/local/bin  
sysconfig: /usr/bin  
Value for scheme.data does not match. Please report this to <<https://github.com/pypa/pip/issues/9617>>  
distutils: /usr/local  
sysconfig: /usr  
Additional context:  
user = False  
home = None  
root = None  
prefix = None  
Removed build tracker: '/tmp/pip-req-tracker-3\_a003f6'

[x] [4] import torch, torchvision  
from torchvision import datasets, models, transforms  
import torch.nn as nn  
import torch.optim as optim  
from torch.optim import lr\_scheduler  
from torch.utils.data import DataLoader  
import time  
from torchsummary import summary  
  
import numpy as np  
import matplotlib.pyplot as plt  
import os  
  
from PIL import Image  
from torch\_lr\_finder import LRFinder

[x] [5] if torch.cuda.is\_available():  
 device = torch.device("cuda:0")  
 print("GPU")  
 torch.cuda.device\_count()  
 torch.cuda.get\_device\_name(0)  
 !nvidia-smi  
else:  
 device = torch.device("cpu")

GPU  
Mon Jan 10 11:33:11 2022  
+-----+  
| NVIDIA-SMI 495.44    Driver Version: 460.32.03    CUDA Version: 11.2 |  
+-----+  
| GPU Name   Persistence-M| Bus-Id   Disp.A | Volatile Uncorr. ECC |  
| Fan Temp Perf Pwr/Usage/Cap| Memory-Usage | GPU-Util Compute M. | MIG M. |  
+-----+  
| 0 Tesla P100-PCIE.. Off | 00000000:00:04.0 off | 0% Default | N/A |  
| N/A 40C P0 27W / 250W | 2MLB / 16280M1B | 0% N/A |  
+-----+  
  
+-----+  
| Processes:                |  
| GPU GI CI    PID Type Process name                          GPU Memory |  
| ID ID                    | Usage |  
+-----+  
| No running processes found |

[x] [6] #parameters  
Batch\_Size = 8  
Number\_of\_Epoch = 30  
Init\_LR\_for\_search = 1e-7

[x] [7] # Applying Transforms to the Data

image\_transforms = [  
 'train': transforms.Compose([  
 transforms.RandomResizedCrop(size = 400, scale=(0.8, 1.0)), #size=256  
 transforms.RandomRotation(degrees=15),  
 transforms.RandomHorizontalFlip(),  
 transforms.CenterCrop(size = 300), #size=224  
 transforms.ToTensor(),  
 transforms.Normalize([0.485, 0.456, 0.406],  
                          [0.229, 0.224, 0.225])

```

    ''',
    'valid': transforms.Compose([
        transforms.Resize(size = 400), #size=256
        transforms.CenterCrop(size = 300), #size=224
        transforms.ToTensor(),
        transforms.Normalize([0.485, 0.456, 0.406],
                           [0.229, 0.224, 0.225])
    ]),
    'test': transforms.Compose([
        transforms.Resize(size = 400), #size=256
        transforms.CenterCrop(size = 300), #size=224
        transforms.ToTensor(),
        transforms.Normalize([0.485, 0.456, 0.406],
                           [0.229, 0.224, 0.225])
    ])
}

# Load the Data
# Set train and valid directory paths
dataset = '/content/gdrive/MyDrive/c2'
#train_path = ("'/content/gdrive/MyDrive/CAD Project Data/challenge1/train")
#val_path = ("'/content/gdrive/MyDrive/CAD Project Data/challenge1/val")

train_directory = os.path.join(dataset, 'train')
valid_directory = os.path.join(dataset, 'val')

# Batch size
bs = Batch Size

# Number of classes
num_classes = len(os.listdir(valid_directory)) #10#2#257
print(num_classes)

# Load Data from folders
data = {
    'train': datasets.ImageFolder(root=train_directory, transform=image_transforms['train']),
    'val': datasets.ImageFolder(root=valid_directory, transform=image_transforms['valid'])
}

# Get a mapping of the indices to the class names, in order to see the output classes of the test images.
idx_to_class = {v: k for k, v in data['train'].class_to_idx.items()}
print(idx_to_class)

# Size of Data, to be used for calculating Average Loss and Accuracy
train_data_size = len(data['train'])
valid_data_size = len(data['val'])

# Create iterators for the Data loaded using DataLoader module
train_data_loader = DataLoader(data['train'], batch_size=bs, shuffle=True)
valid_data_loader = DataLoader(data['val'], batch_size=bs, shuffle=True)

3
{0: 'bcc', 1: 'blk', 2: 'mel'}
```

[9] train\_data\_size, valid\_data\_size

```
(2000, 500)
```

[10] model = models.resnet18(pretrained=True).to(device)

```
model
(Bn2): BatchNorm2d(120, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(downsample): Sequential(
  (0): Conv2d(164, 128, kernel_size=(1, 1), stride=(2, 2), bias=False)
  (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
)
)1: BasicBlock(
  (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace=True)
  (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
)
)layer3: Sequential(
  (0): BasicBlock(
    (conv1): Conv2d(128, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  )
)1: BasicBlock(
  (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace=True)
  (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
)
)layer4: Sequential(
  (0): BasicBlock(
    (conv1): Conv2d(256, 512, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  )
)1: BasicBlock(
  (conv1): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace=True)
  (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
)
)avgpool: AdaptiveAvgPool2d(output_size=(1, 1))
(fc): Linear(in_features=512, out_features=1000, bias=True)
```

[11] # Freeze model parameters

```
for param in model.parameters():
    param.requires_grad = False
```

[12] # Change the final layer of AlexNet Model for Transfer Learning

```
#model.classifier[6] = nn.Linear(4096, num_classes)
model.classifier.add_module("/", nn.LogSoftmax(dim = 1))
num_ftrs = model.fc.in_features
model.fc = nn.Linear(num_ftrs, 3)
model.to(device)
model
```

```
(conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
(Bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(downsample): Sequential(
  (0): Conv2d(164, 128, kernel_size=(1, 1), stride=(2, 2), bias=False)
  (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
)
)1: BasicBlock(
  (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace=True)
  (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
)
)layer3: Sequential(
  (0): BasicBlock(
    (conv1): Conv2d(128, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  )
)1: BasicBlock(
  (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace=True)
  (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
)
)layer4: Sequential(
  (0): BasicBlock(
    (conv1): Conv2d(256, 512, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  )
)1: BasicBlock(
  (conv1): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace=True)
  (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
)
)avgpool: AdaptiveAvgPool2d(output_size=(1, 1))
(fc): Linear(in_features=512, out_features=1000, bias=True)
```

```

        )
    (1: BasicBlock(
        (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
        (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (relu): ReLU(inplace=True)
        (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
        (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
)
(layer4): Sequential(
    (0): BasicBlock(
        (conv1): Conv2d(512, 512, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
        (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (relu): ReLU(inplace=True)
        (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
        (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (downsample): Sequential(
            (0): Conv2d(256, 512, kernel_size=(1, 1), stride=(2, 2), bias=False)
            (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        )
    )
)
(1): BasicBlock(
    (conv1): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
)
)
(avgpool): AdaptiveAvgPool2d(output_size=(1, 1))
(fc): Linear(in_features=512, out_features=3, bias=True)
)

```

✓ [5] summary(model, (3, 224, 224))

	ReLU-21	[-1, 128, 28, 28]	0
Conv2d-22	[-1, 128, 28, 28]	147,456	
BatchNorm2d-23	[-1, 128, 28, 28]	256	
Conv2d-24	[-1, 128, 28, 28]	8,192	
BatchNorm2d-25	[-1, 128, 28, 28]	256	
ReLU-26	[-1, 128, 28, 28]	0	
BasicBlock-27	[-1, 128, 28, 28]	0	
Conv2d-28	[-1, 128, 28, 28]	147,456	
BatchNorm2d-29	[-1, 128, 28, 28]	256	
ReLU-30	[-1, 128, 28, 28]	0	
Conv2d-31	[-1, 128, 28, 28]	147,456	
BatchNorm2d-32	[-1, 128, 28, 28]	256	
ReLU-33	[-1, 128, 28, 28]	0	
BasicBlock-34	[-1, 128, 28, 28]	0	
Conv2d-35	[-1, 256, 14, 14]	294,912	
BatchNorm2d-36	[-1, 256, 14, 14]	512	
ReLU-37	[-1, 256, 14, 14]	0	
Conv2d-38	[-1, 256, 14, 14]	589,824	
BatchNorm2d-39	[-1, 256, 14, 14]	512	
Conv2d-40	[-1, 256, 14, 14]	32,768	
BatchNorm2d-41	[-1, 256, 14, 14]	512	
ReLU-42	[-1, 256, 14, 14]	0	
BasicBlock-43	[-1, 256, 14, 14]	0	
Conv2d-44	[-1, 256, 14, 14]	589,824	
BatchNorm2d-45	[-1, 256, 14, 14]	512	
ReLU-46	[-1, 256, 14, 14]	0	
Conv2d-47	[-1, 256, 14, 14]	589,824	
BatchNorm2d-48	[-1, 256, 14, 14]	512	
ReLU-49	[-1, 256, 14, 14]	0	
BasicBlock-50	[-1, 256, 14, 14]	0	
Conv2d-51	[-1, 512, 7, 7]	1,179,648	
BatchNorm2d-52	[-1, 512, 7, 7]	1,024	
ReLU-53	[-1, 512, 7, 7]	0	
Conv2d-54	[-1, 512, 7, 7]	2,359,296	
BatchNorm2d-55	[-1, 512, 7, 7]	1,024	
Conv2d-56	[-1, 512, 7, 7]	131,072	
BatchNorm2d-57	[-1, 512, 7, 7]	1,024	
ReLU-58	[-1, 512, 7, 7]	0	
BasicBlock-59	[-1, 512, 7, 7]	0	
Conv2d-60	[-1, 512, 7, 7]	2,359,296	
BatchNorm2d-61	[-1, 512, 7, 7]	1,024	
ReLU-62	[-1, 512, 7, 7]	0	
Conv2d-63	[-1, 512, 7, 7]	2,359,296	
BatchNorm2d-64	[-1, 512, 7, 7]	1,024	
ReLU-65	[-1, 512, 7, 7]	0	
BasicBlock-66	[-1, 512, 7, 7]	0	
AdaptiveAvgPool2d-67	[-1, 512, 1, 1]	0	
Linear-68	[-1, 3]	1,539	

=====

Total params: 11,178,051

Trainable params: 1,539

Non-trainable params: 11,176,512

-----

Input size (MB): 0.57

Forward/backward pass size (MB): 62.79

Params size (MB): 42.64

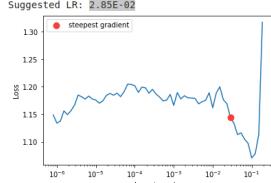
Estimated Total Size (MB): 106.00

-----

```
[14] # Define Optimizer and Loss Function
loss_func = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr = Init_LR_for_search, weight_decay = 1e-2)
lr_finder = LRFinder(model, optimizer, loss_func, device = "cuda:0")
lr_finder.range_test(train_data_loader, end_lr = 100, num_iter = 100)
lr_finder.plot()
#optimizer
```

74% ██████████ 74/100 [01:42<00:31, 1.23s/it]

Stopping early, the loss has diverged
Learning rate search finished. See the graph with {finder\_name}.plot()
LR suggestion: steepest gradient
Suggested LR: 2.85E-02



{<matplotlib.axes.\_subplots.AxesSubplot at 0x7f97362a0650>, 0.028480358684358022}

✓ [15] optimizer = optim.Adam(model.parameters(), lr = 2.85E-02)

✓ [16] decaying\_lr = lr\_scheduler.StepLR(optimizer, step\_size = 4, gamma = 0.1)

✓ [17] def train\_and\_validate(model, loss\_criterion, optimizer, scheduler, epochs):
 ...
 Function to train and validate
 Parameters

```
:param model: Model to train and validate
:param loss_criterion: Loss Criterion to minimize
:param optimizer: Optimizer for computing gradients
:param epochs: Number of epochs (default=25)
```

Returns

```
model: Trained Model with best validation accuracy
history: (dict object): Having training loss, accuracy and validation loss, accuracy
...
```

start = time.time()

history = []

best\_acc = 0.0

for epoch in range(epochs):

epoch\_start = time.time()

print("Epoch: {} / {}".format(epoch+1, epochs))

# Set to training mode

model.train()

# Loss and Accuracy within the epoch

train\_loss = 0.0

train\_acc = 0.0

valid\_loss = 0.0

```

-----  

valid_acc = 0.0  

for i, (inputs, labels) in enumerate(train_data_loader):  

    inputs = inputs.to(device)
    labels = labels.to(device)  

    # Clean existing gradients
    optimizer.zero_grad()  

    # Forward pass - compute outputs on input data using the model
    outputs = model(inputs).to(device)  

    # Compute loss
    loss = loss_criterion(outputs, labels)  

    # Backpropagate the gradients
    loss.backward()  

    # Update the parameters
    optimizer.step()  

    # Compute the total loss for the batch and add it to train_loss
    train_loss += loss.item() * inputs.size(0)  

    # Compute the accuracy
    ret, predictions = torch.max(outputs.data, 1)
    correct_counts = predictions.eq(labels.data.view_as(predictions))  

    # Convert correct_counts to float and then compute the mean
    acc = torch.mean(correct_counts.type(torch.FloatTensor))  

    # Compute total accuracy in the whole batch and add to train_acc
    train_acc += acc.item() * inputs.size(0)  

    #print("Batch number: {:03d}, Training: Loss: {:.4f}, Accuracy: {:.4f}".format(i, loss.item(), acc.item()))
    scheduler.step()  

    # Validation - No gradient tracking needed
    with torch.no_grad():  

        # Set to evaluation mode
        model.eval()  

        # Validation loop
        for j, (inputs, labels) in enumerate(valid_data_loader):  

            inputs = inputs.to(device)
            labels = labels.to(device)  

            # Forward pass - compute outputs on input data using the model
            outputs = model(inputs).to(device)  

            # Compute loss
            loss = loss_criterion(outputs, labels).to(device)  

            # Compute the total loss for the batch and add it to valid_loss
            valid_loss += loss.item() * inputs.size(0)  

            # Calculate validation accuracy
            ret, predictions = torch.max(outputs.data, 1)
            correct_counts = predictions.eq(labels.data.view_as(predictions))  

            # Convert correct_counts to float and then compute the mean
            acc = torch.mean(correct_counts.type(torch.FloatTensor))  

            # Compute total accuracy in the whole batch and add to valid_acc
            valid_acc += acc.item() * inputs.size(0)  

            #print("Validation Batch number: {:03d}, Validation: Loss: {:.4f}, Accuracy: {:.4f}".format(j, loss.item(), acc.item()))  

        # Find average training loss and training accuracy
        avg_train_loss = train_loss/train_data_size
        avg_train_acc = train_acc/train_data_size  

        # Find average training loss and training accuracy
        avg_valid_loss = valid_loss/valid_data_size
        avg_valid_acc = valid_acc/valid_data_size  

        history.append([avg_train_loss, avg_valid_loss, avg_train_acc, avg_valid_acc])
    epoch_end = time.time()  

    print("Epoch : {:03d}, Training: Loss: {:.4f}, Accuracy: {:.4f}%, \n\tValidation : Loss : {:.4f}, Accuracy: {:.4f}%, Time: {:.4f}s".format(epoch+1, avg_train_loss, avg_train_acc*100, avg_valid_loss, avg_valid_acc*100, time.time() - epoch_start))
    # Save if the model has best accuracy till now
    torch.save(model, dataset+'_{model}_{epoch}.pt')

```

```

return model, history

```

```
[18] num_epochs = Number_of_Epoch
trained_model, history = train_and_validate(model, loss_func, optimizer, decaying_lr, num_epochs)
```

```
torch.save(history, dataset+'.history.pt')
```

```
history = np.array(history)
plt.plot(history[:,0:2])
plt.legend(['Tr_Loss', 'Val_Loss'])
plt.xlabel('Epoch Number')
plt.ylabel('Loss')
plt.ylim(0,1)
plt.savefig(dataset+'_loss_curve.png')
plt.show()
```

```

plt.plot(history[:,2:4])
plt.legend(['Tr_Accuracy', 'Val_Accuracy'])
plt.xlabel('Epoch Number')
plt.ylabel('Accuracy')
plt.ylim(0,1)
plt.savefig(dataset+'_accuracy_curve.png')
plt.show()

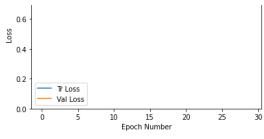
```

```

Epoch : 018, Training: Loss: 0.9044, Accuracy: 73.6500%, Validation : Loss : 0.7759, Accuracy: 79.6000%, Time: 46.8093s
Epoch: 19/30
Epoch : 019, Training: Loss: 0.8947, Accuracy: 73.5000%, Validation : Loss : 0.7924, Accuracy: 80.2000%, Time: 46.9615s
Epoch : 020, Training: Loss: 0.8990, Accuracy: 75.6500%, Validation : Loss : 0.7580, Accuracy: 79.8000%, Time: 46.9416s
Epoch: 21/30
Epoch : 021, Training: Loss: 0.9560, Accuracy: 72.8000%, Validation : Loss : 0.8992, Accuracy: 78.6000%, Time: 47.0617s
Epoch: 22/30
Epoch : 022, Training: Loss: 0.9322, Accuracy: 73.2500%, Validation : Loss : 0.7661, Accuracy: 78.8000%, Time: 46.8114s
Epoch: 23/30
Epoch : 023, Training: Loss: 0.9607, Accuracy: 73.3000%, Validation : Loss : 0.7901, Accuracy: 79.6000%, Time: 46.9446s
Epoch: 24/30
Epoch : 024, Training: Loss: 0.8739, Accuracy: 74.7500%, Validation : Loss : 0.7924, Accuracy: 78.4000%, Time: 47.8557s
Epoch: 25/30
Epoch : 025, Training: Loss: 0.8918, Accuracy: 74.5500%, Validation : Loss : 0.7747, Accuracy: 79.6000%, Time: 48.0402s
Epoch: 26/30
Epoch : 026, Training: Loss: 0.8938, Accuracy: 74.5500%, Validation : Loss : 0.7799, Accuracy: 79.4000%, Time: 46.6367s
Epoch: 27/30
Epoch : 027, Training: Loss: 0.9087, Accuracy: 73.0000%, Validation : Loss : 0.7862, Accuracy: 80.0000%, Time: 45.4838s
Epoch: 28/30
Epoch : 028, Training: Loss: 0.9049, Accuracy: 73.6500%, Validation : Loss : 0.7913, Accuracy: 80.2000%, Time: 45.5587s
Epoch : 29/30
Epoch : 029, Training: Loss: 0.9185, Accuracy: 73.3000%, Validation : Loss : 0.7887, Accuracy: 80.2000%, Time: 45.3660s
Epoch: 30/30
Epoch : 030, Training: Loss: 1.0040, Accuracy: 73.6000%, Validation : Loss : 0.7679, Accuracy: 80.4000%, Time: 45.6235s

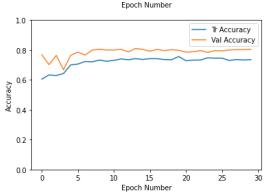
```





```
[19] history = np.array(history)
    plt.plot(history[:,0:2])
    plt.legend(['Tr Loss', 'Val Loss'])
    plt.xlabel('Epoch Number')
    plt.ylabel('Loss')
    plt.ylim(0,1)
    plt.savefig(dataset+'_loss_curve.png')
    plt.show()

    plt.plot(history[:,2:4])
    plt.legend(['Tr Accuracy', 'Val Accuracy'])
    plt.xlabel('Epoch Number')
    plt.ylabel('Accuracy')
    plt.ylim(0,1)
    plt.savefig(dataset+'_accuracy_curve.png')
    plt.show()
```



```
[ ] def predict(model, test_image_name):
    ...
    Function to predict the class of a single test image
    Parameters
        :param model: Model to test
        :param test_image_name: Test image
    ...

    transform = image_transforms['test']

    test_image = Image.open(test_image_name)
    plt.imshow(test_image)

    test_image_tensor = transform(test_image)

    if torch.cuda.is_available():
        test_image_tensor = test_image_tensor.view(1, 3, 224, 224).cuda()
    else:
        test_image_tensor = test_image_tensor.view(1, 3, 224, 224)

    with torch.no_grad():
        model.eval()
        # Model outputs log probabilities
        out = model(test_image_tensor)
        ps = torch.exp(out)
        topk, topclass = ps.topk(3, dim=1)
        for i in range(3):
            print("Prediction", i+1, ":", idx_to_class[topclass.numpy()[0][i]], ", Score: ", topk.numpy()[0][i])

[ ] # predict(trained_model, 'caltec256subset/test/009_0098.jpg')
# predict(trained_model, 'caltec256subset/test/090_0107.jpg')
# predict(trained_model, 'caltec256subset/test/150_0044.jpg')
```

✓ 1s completed at 1:05 PM

